

## Kubernetes, Daemomset, taints and toleration (ASSIGNMENT-13)

### 1. AWS Setup for EKS Cluster

Step 1: Configure AWS CLI

*Command: aws configure*

---

Inputs:

- AWS Access Key ID
- AWS Secret Access Key
- Region (e.g., ap-south-1)
- Output format (e.g., json)

Step 2: Verify Identity

*Command: aws sts get-caller-identity*

---

Returns your AWS Account ID, User ARN, and User ID.

```
PS C:\WINDOWS\system32> aws configure
AWS Access Key ID [*****ADWB]: AKIAT5VFFI0DHZFLBIF7
AWS Secret Access Key [*****rplwb]: X2oJtA/cErSGclmsa2/Vz5MvqcCp+Ltv9Mtt+ogE
Default region name [us-east-1]: ap-south-1
Default output format [json]:
PS C:\WINDOWS\system32> aws sts get-caller-identity
{
  "UserId": "269855572870",
  "Account": "269855572870",
  "Arn": "arn:aws:iam::269855572870:root"
}
```

### 2. Create EKS Cluster Using eksctl

*Command:*

---

eksctl create cluster --name demo-cluster --region ap-south-1 --nodegroup-name standard-workers --node-type t3.medium --nodes 2 --nodes-min 1 --nodes-max 3 --managed

```

PS C:\WINDOWS\system32> eksctl create cluster --region ap-south-1 --name testing-cluster --node-type t3.medium --nodes 2 --nodes-min 2 --
nodes-max 3 --managed
2025-06-26 12:25:29 [i] eksctl version 0.210.0
2025-06-26 12:25:29 [i] using region ap-south-1
2025-06-26 12:25:30 [i] setting availability zones to [ap-south-1a ap-south-1c ap-south-1b]
2025-06-26 12:25:30 [i] subnets for ap-south-1a - public:192.168.0.0/19 private:192.168.96.0/19
2025-06-26 12:25:30 [i] subnets for ap-south-1c - public:192.168.32.0/19 private:192.168.128.0/19
2025-06-26 12:25:30 [i] subnets for ap-south-1b - public:192.168.64.0/19 private:192.168.160.0/19
2025-06-26 12:25:30 [i] nodegroup "ng-1ba28734" will use "" [AmazonLinux2023/1.32]
2025-06-26 12:25:30 [i] using Kubernetes version 1.32
2025-06-26 12:25:30 [i] creating EKS cluster "testing-cluster" in "ap-south-1" region with managed nodes
2025-06-26 12:25:30 [i] will create 2 separate CloudFormation stacks for cluster itself and the initial managed nodegroup
2025-06-26 12:25:30 [i] if you encounter any issues, check CloudFormation console or try 'eksctl utils describe-stacks --region=ap-south
-1 --cluster=testing-cluster'
2025-06-26 12:25:30 [i] Kubernetes API endpoint access will use default of {publicAccess=true, privateAccess=false} for cluster "testing
-cluster" in "ap-south-1"
2025-06-26 12:25:30 [i] CloudWatch logging will not be enabled for cluster "testing-cluster" in "ap-south-1"
2025-06-26 12:25:30 [i] you can enable it with 'eksctl utils update-cluster-logging --enable-types={SPECIFY-YOUR-LOG-TYPES-HERE (e.g. al
l)} --region=ap-south-1 --cluster=testing-cluster'
2025-06-26 12:25:30 [i] default addons metrics-server, vpc-cni, kube-proxy, coredns were not specified, will install them as EKS addons
2025-06-26 12:25:30 [i]
2 sequential tasks: { create cluster control plane "testing-cluster",
  2 sequential sub-tasks: {
    2 sequential sub-tasks: {
      1 task: { create addons },
      wait for control plane to become ready,
    },
    create managed nodegroup "ng-1ba28734",
  },
}
2025-06-26 12:25:30 [i] building cluster stack "eksctl-testing-cluster-cluster"
2025-06-26 12:25:31 [i] deploying stack "eksctl-testing-cluster-cluster"
2025-06-26 12:33:35 [i] successfully created addon: metrics-server
2025-06-26 12:33:35 [i] recommended policies were found for "vpc-cni" addon, but since OIDC is disabled on the cluster, eksctl
cannot configure the requested permissions; the recommended way to provide IAM permissions for "vpc-cni" addon is via pod identi
ty associations; after addon creation is completed, add all recommended policies to the config file, under 'addon.PodIdentityAss
ociations', and run 'eksctl update addon'
2025-06-26 12:33:35 [i] creating addon: vpc-cni
2025-06-26 12:33:36 [i] successfully created addon: vpc-cni
2025-06-26 12:33:36 [i] creating addon: kube-proxy
2025-06-26 12:33:37 [i] successfully created addon: kube-proxy
2025-06-26 12:33:37 [i] creating addon: coredns
2025-06-26 12:33:38 [i] successfully created addon: coredns
2025-06-26 12:35:39 [i] building managed nodegroup stack "eksctl-testing-cluster-nodegroup-ng-1ba28734"
2025-06-26 12:41:02 [i] deploying stack "eksctl-testing-cluster-nodegroup-ng-1ba28734"
2025-06-26 12:41:03 [i] waiting for CloudFormation stack "eksctl-testing-cluster-nodegroup-ng-1ba28734"
2025-06-26 12:41:33 [i] waiting for CloudFormation stack "eksctl-testing-cluster-nodegroup-ng-1ba28734"
2025-06-26 12:42:09 [i] waiting for CloudFormation stack "eksctl-testing-cluster-nodegroup-ng-1ba28734"
2025-06-26 12:43:46 [i] waiting for CloudFormation stack "eksctl-testing-cluster-nodegroup-ng-1ba28734"
2025-06-26 12:43:46 [i] waiting for the control plane to become ready
2025-06-26 12:43:48 [i] saved kubeconfig as "C:\\Users\\Ayush Singh\\.kube\\config"
2025-06-26 12:43:48 [i] no tasks
2025-06-26 12:43:48 [i] all EKS cluster resources for "testing-cluster" have been created
2025-06-26 12:43:48 [i] nodegroup "ng-1ba28734" has 2 node(s)
2025-06-26 12:43:48 [i] node "ip-192-168-41-155.ap-south-1.compute.internal" is ready
2025-06-26 12:43:48 [i] node "ip-192-168-64-17.ap-south-1.compute.internal" is ready
2025-06-26 12:43:48 [i] waiting for at least 2 node(s) to become ready in "ng-1ba28734"
2025-06-26 12:43:48 [i] nodegroup "ng-1ba28734" has 2 node(s)
2025-06-26 12:43:48 [i] node "ip-192-168-41-155.ap-south-1.compute.internal" is ready
2025-06-26 12:43:48 [i] node "ip-192-168-64-17.ap-south-1.compute.internal" is ready
2025-06-26 12:43:48 [i] created 1 managed nodegroup(s) in cluster "testing-cluster"
2025-06-26 12:43:50 [i] kubectl command should work with "C:\\Users\\Ayush Singh\\.kube\\config", try 'kubectl get nodes'
2025-06-26 12:43:50 [i] EKS cluster "testing-cluster" in "ap-south-1" region is ready

```

### 3. Kubernetes Routing Structure

Manual EC2 Hosting:

EC2 → Application Host → Target Group → Load Balancer → DNS (Access)

Kubernetes Routing:

Service → Deployment → Pods

```

PS C:\WINDOWS\system32> aws eks update-kubeconfig --region ap-south-1 --name testing-cluster
Added new context arn:aws:eks:ap-south-1:269855572870:cluster/testing-cluster to C:\Users\Ayush Singh\.kube\config

```

*Apply Commands:*

---

```
kubectl apply -f deployment.yaml
kubectl apply -f service.yaml
```

#### 4. Clone GitHub Manifests Repository

Repository: <https://github.com/sibasish934/kubernetes-manifests>

##### *Commands:*

---

```
git clone https://github.com/sibasish934/kubernetes-manifests.git
cd kubernetes-manifests
```

```
PS C:\WINDOWS\system32> git clone https://github.com/sibasish934/kubernetes-manifests.git
Cloning into 'kubernetes-manifests'...
remote: Enumerating objects: 9, done.
remote: Counting objects: 100% (9/9), done.
remote: Compressing objects: 100% (8/8), done.
remote: Total 9 (delta 1), reused 9 (delta 1), pack-reused 0 (from 0)
Receiving objects: 100% (9/9), done.
Resolving deltas: 100% (1/1), done.
PS C:\WINDOWS\system32> cd kubernetes-manifests
```

```
PS C:\WINDOWS\system32\kubernetes-manifests> kubectl apply -f deployment.yaml
deployment.apps/devops-deployment created
PS C:\WINDOWS\system32\kubernetes-manifests> kubectl apply -f service.yaml
service/nginx-service created
PS C:\WINDOWS\system32\kubernetes-manifests> kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
devops-deployment-96b9d695-nms4v    1/1     Running   0           13s
```

#### 5. Resource Requests and Limits

Problem: One pod can consume too many resources and affect others.

Solution: Use resource requests and limits.

##### *resources:*

###### *requests:*

*cpu: "250m"*

*memory: "64Mi"*

###### *limits:*

*cpu: "500m"*

*memory: "128Mi"*

*Apply: kubectl apply -f resources-demo.yaml*

---

```
PS C:\WINDOWS\system32\kubernetes-manifests> notepad resources-demo.yaml
PS C:\WINDOWS\system32\kubernetes-manifests> kubectl apply -f resources-demo.yaml -n devops
pod/resource-check created
PS C:\WINDOWS\system32\kubernetes-manifests> kubectl get pod resource-check -n devops
NAME          READY   STATUS    RESTARTS   AGE
resource-check 1/1     Running   0          9s
PS C:\WINDOWS\system32\kubernetes-manifests> kubectl describe pod resource-check -n devops
Name:          resource-check
Namespace:     devops
Priority:       0
Service Account: default
Node:          ip-192-168-41-155.ap-south-1.compute.internal/192.168.41.155
Start Time:    Thu, 26 Jun 2025 13:39:30 +0530
Labels:        <none>
Annotations:   <none>
Status:        Running
IP:            192.168.58.2
IPs:
  IP: 192.168.58.2
Containers:
  nginx:
    Container ID:  containerd://a831b9cf5d6d383ac0a2694197fbb6df5684a5493a22c04cc05cdbedca3b8bcb
    Image:         nginx
    Image ID:      docker.io/library/nginx@sha256:dc53c8f25a10f9109190ed5b59bda2d707a3bde0e45857ce9e1efaa32ff9cbc1
    Port:          <none>
    Host Port:     <none>
    State:         Running
      Started:     Thu, 26 Jun 2025 13:39:32 +0530
    Ready:         True
    Restart Count: 0
    Limits:
      cpu:         500m
      memory:      128Mi
    Requests:
      cpu:         250m
      memory:      64Mi
    Environment:   <none>
```

## 6. Deploy a DaemonSet

Purpose: Ensure that a pod runs on every node.

Use-case: Logging agents, monitoring agents.

*Command: `kubectl apply -f daemonset.yaml`*

Instances (2) Info

Last updated less than a minute ago

Connect

Instance state

Actions

Launch instances

All states

☐

testing-cluster...

i-0f9197f8ad938d97a

Running

t3.medium

3/3 checks passed

View alarms +

ap-south-1c

ec2-3-1

☐

testing-cluster...

i-0e47759e5f3e93f34

Running

t3.medium

3/3 checks passed

View alarms +

ap-south-1b

ec2-13-

```
PS C:\WINDOWS\system32\kubernetes-manifests> kubectl apply -f daemonset.yaml -n devops
daemonset.apps/nginx-daemonset created
PS C:\WINDOWS\system32\kubernetes-manifests> kubectl get pods -n devops -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE
devops-deployment-96b9d695-nf2gc    1/1     Running   0          22m   192.168.83.118  ip-192-168-64-17.ap-south-1.compute.internal
nginx-daemonset-ck9tl                1/1     Running   0          2m54s 192.168.35.206  ip-192-168-41-155.ap-south-1.compute.internal
nginx-daemonset-tbm8l                1/1     Running   0          2m54s 192.168.85.250  ip-192-168-64-17.ap-south-1.compute.internal
```

## 7. Taints and Tolerations

Use taints to prevent pods from being scheduled unless they have matching tolerations.

*Taint a node:*

`kubectl taint nodes <node-name> key=value:NoSchedule`

### *Toleration YAML:*

---

tolerations:

- key: "key"
- operator: "Equal"
- value: "value"
- effect: "NoSchedule"

### *Apply: kubectl apply -f toleration-demo.yaml*

---

```
PS C:\WINDOWS\system32\kubernetes-manifests> notepad toleration-demo.yaml
PS C:\WINDOWS\system32\kubernetes-manifests> kubectl apply -f toleration-demo.yaml -n devops
pod/tolerant-pod created
PS C:\WINDOWS\system32\kubernetes-manifests> kubectl get pods -n devops -o wide
>>


| NAME                             | NOMINATED NODE | READINESS GATES | READY | STATUS  | RESTARTS | AGE   | IP             | NODE                                          |
|----------------------------------|----------------|-----------------|-------|---------|----------|-------|----------------|-----------------------------------------------|
| devops-deployment-96b9d695-nf2gc | <none>         | <none>          | 1/1   | Running | 0        | 37m   | 192.168.83.118 | ip-192-168-64-17.ap-south-1.compute.internal  |
| nginx-daemonset-ck9tl            | <none>         | <none>          | 1/1   | Running | 0        | 17m   | 192.168.35.206 | ip-192-168-41-155.ap-south-1.compute.internal |
| nginx-daemonset-tbm8l            | <none>         | <none>          | 1/1   | Running | 0        | 17m   | 192.168.85.250 | ip-192-168-64-17.ap-south-1.compute.internal  |
| resource-check                   | <none>         | <none>          | 1/1   | Running | 0        | 5m25s | 192.168.58.2   | ip-192-168-41-155.ap-south-1.compute.internal |
| tolerant-pod                     | <none>         | <none>          | 1/1   | Running | 0        | 24s   | 192.168.84.37  | ip-192-168-64-17.ap-south-1.compute.internal  |


```

## 8. Expose App to Internet (optional)

Expose deployment via LoadBalancer service:

```
kubectl expose deployment <deployment-name> --  
type=LoadBalancer --name=<service-name>
```

---

Check external IP:

```
kubectl get svc
```

---

## 9. Useful kubectl Commands

```
kubectl get all                   # View all resources
```

```
kubectl get nodes               # List cluster nodes
```

```
kubectl get pods -o wide       # View pod info with node IPs
```

```
kubectl describe pod <name>     # Describe pod in detail
```

```
kubectl logs <pod-name>       # View logs
```

```
kubectl exec -it <pod> -- sh    # Shell access to pod
```

```
kubectrl delete -f <file.yaml> # Delete a resource
```

---

## 10. Clean Up Resources

Delete all resources in 'devops' namespace:

```
kubectrl delete all --all -n devops
```

---

Delete individual pods if needed:

```
kubectrl delete pod resource-check -n devops  
kubectrl delete pod tolerant-pod -n devops
```

---

Delete the namespace (optional):

```
kubectrl delete namespace devops
```

---

Delete the EKS cluster (IMPORTANT):

```
eksctl delete cluster --name demo-cluster --region ap-south-1
```

---