

Kubernet Application & Archicecture(components) [Assignment-12]

1. Container Management Tool: Docker Swarm

- Docker Swarm is Docker's native clustering and container orchestration tool.
- It allows users to manage multiple Docker containers deployed across multiple host machines.

Disadvantages:

- Manual auto-scaling: No built-in support for dynamic scaling of containers.
 - Manual load balancing: Requires manual configuration for traffic distribution.
-

2. Evolution of Kubernetes (K8s)

- Google Borg: Internal container management system at Google.
 - Kubernetes was developed by Google using Borg's principles and donated to CNCF (Cloud Native Computing Foundation) in 2014.
 - Open-source and No Vendor Lock-in: Unlike proprietary tools, Kubernetes is not tied to a single cloud provider.
-

3. Kubernetes Architecture and Cluster Components

Control Plane Components:

1. etcd

- Key-value store that stores all cluster data, configuration, and state.

2. kube-apiserver

- Central management entity that exposes the Kubernetes API.
- Handles communication between components.

3. scheduler

- Assigns newly created Pods to nodes based on resource requirements and availability.

4. controller manager

- Manages different controllers (e.g., replication, endpoint, namespace).
- Ensures desired state of the cluster is maintained.

5. cloud-controller-manager (CCM)

- Integrates Kubernetes with cloud providers like AWS, GCP, etc.
- Manages cloud-specific services (load balancers, node information).

Node Components (Worker Nodes):

1. kubelet

- Agent that runs on each node and communicates with the control plane.
- Ensures containers are running in Pods as expected.

2. kube-proxy

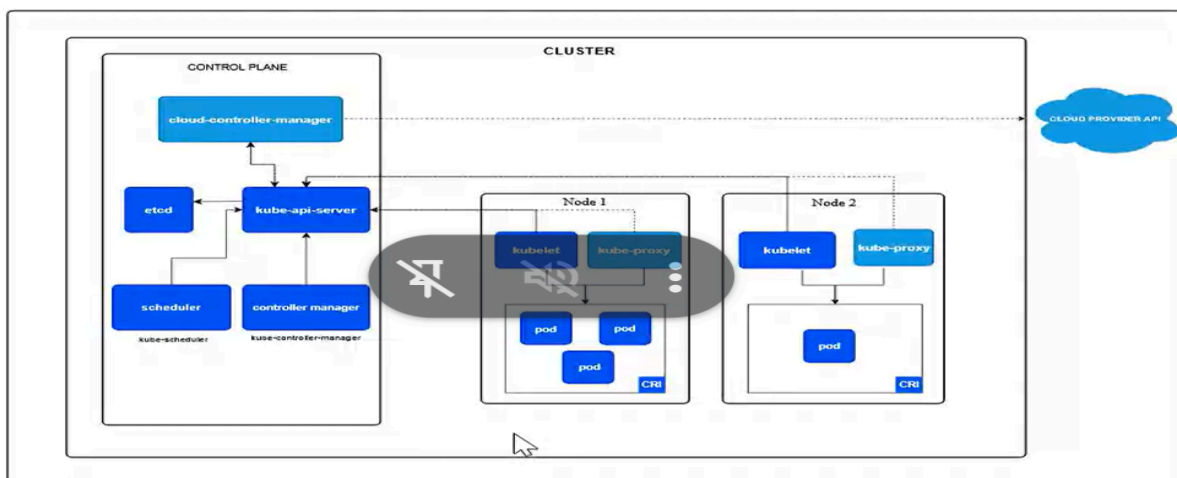
- Handles networking and load-balancing between Pods.
- Manages network rules on nodes.

3. Pod

- Smallest deployable unit in Kubernetes.
- Can contain one or more containers sharing the same network and storage.

4. CRI (Container Runtime Interface)

- Allows Kubernetes to use different container runtimes (e.g., containerd, CRI-O).
- Responsible for pulling images and starting containers.



Kubernetes Cluster Workflow to Deploy 2 NGINX Containers

Step 1: Write a Deployment YAML File

Create a file named `nginx-deployment.yaml`:

```
apiVersion: apps/v1

kind: Deployment

metadata:
  name: nginx-deployment

spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
            - containerPort: 80
```

Step 2: Apply the Deployment

Run the following command to create the deployment:

```
kubectl apply -f nginx-deployment.yaml
```

Step 3: Kubernetes Control Plane Processes the Request

- **kube-apiserver** receives the deployment request.
- The request is stored in **etcd** (cluster state database).
- **Controller Manager** notices the desired state (2 replicas) and creates 2 Pods.
- **Scheduler** assigns the 2 Pods to available worker nodes.

Step 4: Worker Node Executes Pods

- **kubelet** on the assigned node pulls the `nginx:latest` image.
- It launches containers using the container runtime (**containerd/CRI**).
- The containers run inside Pods.

Step 5: kube-proxy Manages Networking

- **kube-proxy** sets up necessary network rules to route traffic to the Pods.
- If you expose the deployment using a **Service**, kube-proxy routes incoming traffic accordingly.

(Optional) Step 6: Expose the Deployment

To make NGINX accessible outside the cluster:

```
kubectl expose deployment nginx-deployment --port=80  
--type=NodePort
```

Verification

```
kubectl get pods
```

```
kubectl get deployments
```

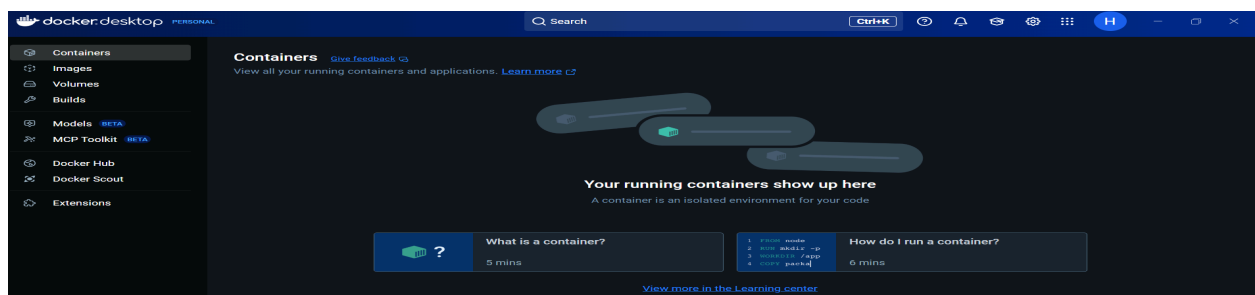
```
kubectl get services
```

Lab Session commands: -

1. Create the Kind Cluster in Local.
 - a. Install Kind (use choco to install)

```
PS C:\WINDOWS\system32> choco install kind  
Chocolatey v2.3.0  
Installing the following packages:  
kind  
By installing, you accept licenses for the packages.  
Downloading package from source 'https://community.chocolatey.org/api/v2/'  
Progress: Downloading kind 0.29.0... 100%  
kind v0.29.0 [Approved]  
kind package files install completed. Performing other installation steps.  
The package kind wants to run 'chocolateyinstall.ps1'.  
Note: If you don't run this script, the installation will fail.  
Note: To confirm automatically next time, use '-y' or consider:  
choco feature enable --n allowGlobalConfirmation  
Do you want to run the script?([Y]es/[A]ll - yes to all/[N]o/[P]rint): A  
Downloading kind 64 bit  
from 'https://github.com/kubernetes-sigs/kind/releases/download/v0.29.0/kind-windows-amd64'.  
Progress: 100% - Completed download of C:\ProgramData\chocolatey\lib\kind\kind.exe (10.74 MB).  
Download of kind.exe (10.74 MB) completed.  
Hashes match.  
ShimGen has successfully created a shim for kind.exe  
The install of kind was successful.  
Software install location not explicitly set, it could be in package or  
default install location of installer.  
Chocolatey installed 1/1 packages.  
See the log for details (C:\ProgramData\chocolatey\logs\chocolatey.log).
```

- b. Install docker desktop



- c. Use the command: - kind create cluster --name testing-cluster
--config cluster.yaml

Cluster.yaml file:-

kind: Cluster

apiVersion: kind.x-k8s.io/v1alpha4

nodes:

- role: control-plane

image:

kindest/node:v1.30.0@sha256:047357ac0cfea04663786a612ba1eaba9
702bef25227a794b52890dd8bcd692e

- role: worker

image:

kindest/node:v1.29.4@sha256:3abb816a5b1061fb15c6e9e60856ec40d
56b7b52bcea5f5f1350bc6e2320b6f8

Images

[Give feedback](#)

View and manage your local and Docker Hub images. [Learn more](#)

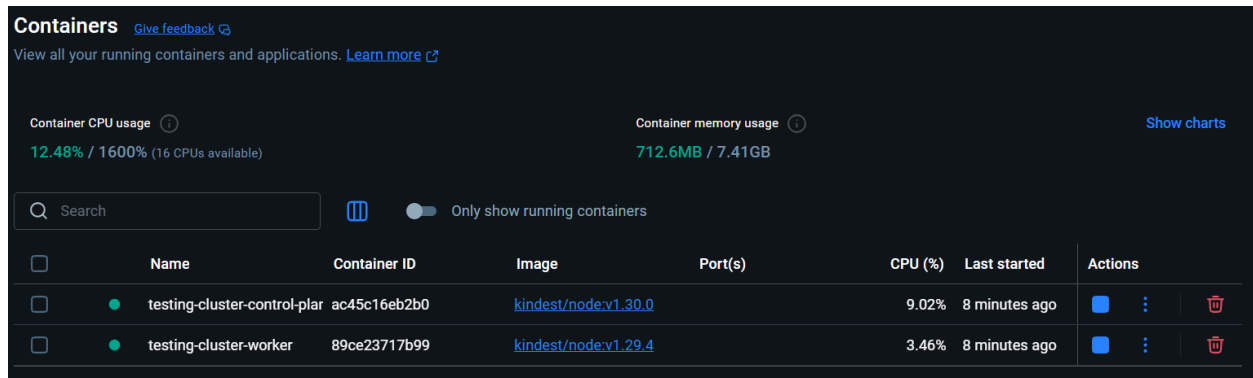
Local

My Hub

1.97 GB / 1.77 GB in use 2 images

Last refresh: 11 minutes ago

<input type="checkbox"/>	Name	Tag	Image ID	Created	Size	Actions
<input type="checkbox"/>	kindest/node	<none>	3abb816a5b10	1 year ago	1.39 GB	
<input type="checkbox"/>	kindest/node	<none>	047357ac0cfe	1 year ago	1.41 GB	



2. Or You can use online playground at killercode.com
3. Use these commands and take screenshot of the outputs as the assignment.
 - a. `Kubectl get ns`

```
PS C:\Users\Ayush Singh\Downloads\DevOps-Assignment> kubectl get ns
NAME                STATUS    AGE
default             Active    12m
kube-node-lease     Active    12m
kube-public         Active    12m
kube-system         Active    12m
local-path-storage  Active    12m
```

- b. `Kubectl get pods`
- c. `Kubectl get pods -n <namespace-name>`
- d. `Kubectl describe pod/<pod_name>`

```
PS C:\Users\Ayush Singh\Downloads\DevOps-Assignment> kubectl get pods
NAME        READY   STATUS    RESTARTS   AGE
devops-pod  1/1     Running   0           2m33s
```

- e. `Kubectl logs pod/<pod_name>`


```

PS C:\Users\sidd2\Downloads\kubernetes_class> kubectl describe pod devops-pod
Name:          devops-pod
Namespace:     default
Priority:       0
Service Account: default
Node:          testing-cluster-worker/172.18.0.2
Start Time:    Wed, 25 Jun 2025 16:20:15 +0530
Labels:        run=devops-pod
Annotations:   <none>
Status:        Running
IP:            10.244.1.2
IPs:
  IP: 10.244.1.2
Containers:
  devops-pod:
    Container ID:  containerd://ccd0f4f680f26a61d919766850e24c192955f4ad35e18393a398810ad203a8ef
    Image:         nginx
    Image ID:      docker.io/library/nginx@sha256:dc53c8f25a10f9109190ed5b59bda2d707a3bde0e45857ce9e1e
    Port:          <none>
    Host Port:     <none>
    State:         Running
      Started:     Wed, 25 Jun 2025 16:21:03 +0530
    Ready:         True
    Restart Count:  0
    Environment:   <none>
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-q4x2n (ro)
Conditions:
  Type                               Status
  PodReadyToStartContainers         True
  Initialized                       True
  Ready                             True
  ContainersReady                   True
  PodScheduled                      True
Volumes:
  kube-api-access-q4x2n:
    Type:          Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName:  kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI:    true
QoS Class:         BestEffort
Node-Selectors:    <none>
Tolerations:       node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
                   node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type    Reason      Age   From          Message
  ----    -

```

f. Kubectl run devops-pod --image= nginx

```

PS C:\Users\Ayush Singh\Downloads\DevOps-Assignment> kubectl run devops-pod --image=nginx
pod/devops-pod created

```

g. Kubectl get pods -n kube-system :- to display the components of the control plane pod of Kubernetes.

```
PS C:\Users\Ayush Singh\Downloads\DevOps-Assignment> kubectl get pods -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
coredns-7db6d8ff4d-kk6cx	1/1	Running	0	17m
coredns-7db6d8ff4d-wwjxm	1/1	Running	0	17m
etcd-testing-cluster-control-plane	1/1	Running	0	17m
kindnet-q71zq	1/1	Running	0	17m
kindnet-sjknx	1/1	Running	0	17m
kube-apiserver-testing-cluster-control-plane	1/1	Running	0	17m
kube-controller-manager-testing-cluster-control-plane	1/1	Running	0	17m
kube-proxy-6nhw	1/1	Running	0	17m
kube-proxy-vslth	1/1	Running	0	17m
kube-scheduler-testing-cluster-control-plane	1/1	Running	0	17m

4. Create EKS cluster in the AWS.

a. Command to create it:- eksctl create cluster \

--name my-eks-cluster \

--region us-west-2 \

--nodegroup-name linux-nodes \

--node-type t3.medium \

--nodes 2 \

--nodes-min 1 \

--nodes-max 3 \

--managed è check the cloudformation template and analyze it what is being created.

```

PS C:\Users\sidd2\Downloads\kubernetes_class> kubectl logs devops-pod
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2025/06/25 10:51:04 [notice] 1#1: using the "epoll" event method
2025/06/25 10:51:04 [notice] 1#1: nginx/1.29.0
2025/06/25 10:51:04 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14+deb12u1)
2025/06/25 10:51:04 [notice] 1#1: OS: Linux 5.15.167.4-microsoft-standard-WSL2
2025/06/25 10:51:04 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2025/06/25 10:51:04 [notice] 1#1: start worker processes
2025/06/25 10:51:04 [notice] 1#1: start worker process 32
2025/06/25 10:51:04 [notice] 1#1: start worker process 33
2025/06/25 10:51:04 [notice] 1#1: start worker process 34
2025/06/25 10:51:04 [notice] 1#1: start worker process 35
2025/06/25 10:51:04 [notice] 1#1: start worker process 36
2025/06/25 10:51:04 [notice] 1#1: start worker process 37
2025/06/25 10:51:04 [notice] 1#1: start worker process 38
2025/06/25 10:51:04 [notice] 1#1: start worker process 39
PS C:\Users\sidd2\Downloads\kubernetes_class>

```

b. Command to delete the Resource: - `eksctl delete cluster --region <region-code> --name <name of the cluster.>`

c. Also execute the sa

me commands in the EKS cluster also

d. Create two namespaces in EKS cluster and then create two pod of apache webserver (httpd) in that particular namespaces only,

```

PS C:\Users\sidd2\Downloads\kubernetes_class> kubectl create ns team-a
namespace/team-a created
PS C:\Users\sidd2\Downloads\kubernetes_class> kubectl create ns team-b
namespace/team-b created

```

```

Error: failed to create cluster "my-eks-cluster"
PS C:\Users\sidd2\Downloads\kubernetes_class> kubectl create ns team-a
namespace/team-a created
PS C:\Users\sidd2\Downloads\kubernetes_class> kubectl create ns team-b
namespace/team-b created
PS C:\Users\sidd2\Downloads\kubernetes_class>
PS C:\Users\sidd2\Downloads\kubernetes_class> kubectl run apache1 --image=httpd -n team-a
pod/apache1 created
PS C:\Users\sidd2\Downloads\kubernetes_class> kubectl run apache2 --image=httpd -n team-a
pod/apache2 created
PS C:\Users\sidd2\Downloads\kubernetes_class>
PS C:\Users\sidd2\Downloads\kubernetes_class> kubectl run apache1 --image=httpd -n team-b
pod/apache1 created
PS C:\Users\sidd2\Downloads\kubernetes_class> kubectl run apache2 --image=httpd -n team-b
pod/apache2 created

```